

Various CNN Networks On MNIST Data Set :

In [2]:

```
from __future__ import print_function
from datetime import datetime
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
from keras.initializers import he_normal
from keras.layers.normalization import BatchNormalization
batch_size = 128
num_classes = 10
epochs = 12
# input image dimensions
img_rows, img_cols = 28, 28
```

Using TensorFlow backend.

In [4]:

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

In [5]:

```
if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

In [6]:

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty):
    fig = plt.figure( facecolor='y', edgecolor='k')
    plt.plot(x, vy, 'b', label="Validation Loss")
    plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel('Epochs')
    plt.ylabel('Categorical Crossentropy Loss')
```

```
plt.legend()
plt.grid()
plt.show()
```

1st Model With 3 ConvNet & 3*3 Kernel Size :

In [7]:

```
convnet=Sequential() # Initializing the model
convnet.add(Conv2D(32,kernel_size=(3,3),activation='relu',input_shape=input_shape))
convnet.add(Conv2D(64,kernel_size=(3,3),activation='relu'))
convnet.add(Dropout(0.25))
convnet.add(Conv2D(128,kernel_size=(3,3),activation='relu'))
#maxpooling by (2,2) ,dropout,flattening
convnet.add(MaxPooling2D(pool_size=(2,2)))
convnet.add(Dropout(0.25))
convnet.add(Flatten())
#hidden_layer
convnet.add(Dense(256,activation='relu',kernel_initializer=he_normal(seed=None)))
convnet.add(Dropout(0.5))
convnet.add(Dense(num_classes,activation='softmax'))
print(convnet.summary())
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18496
dropout_1 (Dropout)	(None, 24, 24, 64)	0
conv2d_3 (Conv2D)	(None, 22, 22, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 11, 11, 128)	0
dropout_2 (Dropout)	(None, 11, 11, 128)	0
flatten_1 (Flatten)	(None, 15488)	0
dense_1 (Dense)	(None, 256)	3965184
dropout_3 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2570
Total params: 4,060,426		
Trainable params: 4,060,426		
Non-trainable params: 0		

None

In [8]:

```
convnet.compile(optimizer=keras.optimizers.Adam(),loss=keras.losses.categorical_crossentropy,metrics=['accuracy'])
convnet_history=convnet.fit(x_train,y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 231s 4ms/step - loss: 0.1775 - accuracy: 0.9453 - val_loss: 0.0435 - val_accuracy: 0.9857

Epoch 2/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0601 - accuracy: 0.9823 - val_loss: 0.0385 - val_accuracy: 0.9874

Epoch 3/12

60000/60000 [=====] - 226s 4ms/step - loss: 0.0467 - accuracy: 0.9857 - val_loss: 0.0335 - val_accuracy: 0.9892

Epoch 4/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0365 - accuracy: 0.9886 - val_loss: 0.0285 - val_accuracy: 0.9910

```

al_loss: 0.0235 - val_accuracy: 0.9918
Epoch 5/12
60000/60000 [=====] - 237s 4ms/step - loss: 0.0288 - accuracy: 0.9912 - v
al_loss: 0.0224 - val_accuracy: 0.9924
Epoch 6/12
60000/60000 [=====] - 237s 4ms/step - loss: 0.0241 - accuracy: 0.9923 - v
al_loss: 0.0224 - val_accuracy: 0.9929
Epoch 7/12
60000/60000 [=====] - 231s 4ms/step - loss: 0.0221 - accuracy: 0.9929 - v
al_loss: 0.0254 - val_accuracy: 0.9927
Epoch 8/12
60000/60000 [=====] - 225s 4ms/step - loss: 0.0212 - accuracy: 0.9932 - v
al_loss: 0.0258 - val_accuracy: 0.9927
Epoch 9/12
60000/60000 [=====] - 230s 4ms/step - loss: 0.0184 - accuracy: 0.9944 - v
al_loss: 0.0222 - val_accuracy: 0.9936
Epoch 10/12
60000/60000 [=====] - 233s 4ms/step - loss: 0.0152 - accuracy: 0.9951 - v
al_loss: 0.0241 - val_accuracy: 0.9936
Epoch 11/12
60000/60000 [=====] - 222s 4ms/step - loss: 0.0151 - accuracy: 0.9952 - v
al_loss: 0.0253 - val_accuracy: 0.9922
Epoch 12/12
60000/60000 [=====] - 222s 4ms/step - loss: 0.0145 - accuracy: 0.9953 - v
al_loss: 0.0286 - val_accuracy: 0.9926

```

In [10]:

```


score=convnet.evaluate(x_test,y_test,verbose=0)
test_score=score[0]
test_accuracy=score[1]
train_accuracy=max(convnet_history.history['accuracy'])
print('test score :',test_score)
print('test sccuracy :',test_accuracy)
# error plot
x=list(range(1,epochs+1))
vy=convnet_history.history['val_loss'] #validation loss
ty=convnet_history.history['loss'] # train loss
plt_dynamic(x, vy, ty)

```

```

test score : 0.028591023394818865
test sccuracy : 0.9926000237464905

```



2nd Model CNN with 5 ConvNet & 5*5 Kernel Size :

In [11]:

```

convnet2=Sequential() # Initializing the model
# First ConvNet
convnet2.add(Conv2D(32,kernel_size=(5,5),activation='relu',padding='same',input_shape=input_shape))
convnet2.add(Conv2D(64,kernel_size=(5,5),padding='same',activation='relu')) #Second Convnet
convnet2.add(MaxPooling2D(pool_size=(2,2)))
convnet2.add(Dropout(0.25))
convnet2.add(Conv2D(96,kernel_size=(5,5),padding='same',activation='relu')) # 3rd ConvNet
#maxpooling by (2,2) ,dropout,flattening
convnet2.add(MaxPooling2D(pool_size=(2,2)))

```

```

convnet2.add(Dropout(0.25))
convnet2.add(Conv2D(128, kernel_size=(5,5), padding='same', activation='relu')) #fourth Convnet
convnet2.add(MaxPooling2D(pool_size=(2,2)))
convnet2.add(Dropout(0.25))
convnet2.add(Conv2D(164, kernel_size=(5,5), padding='same', activation='relu')) #fifth Convnet
convnet2.add(MaxPooling2D(pool_size=(2,2)))
convnet2.add(Dropout(0.25))
convnet2.add(Flatten())
#hidden_layer
convnet2.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
convnet2.add(BatchNormalization())
convnet2.add(Dropout(0.5))
convnet2.add(Dense(num_classes, activation='softmax'))
print(convnet2.summary())

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 28, 28, 32)	832
conv2d_5 (Conv2D)	(None, 28, 28, 64)	51264
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_4 (Dropout)	(None, 14, 14, 64)	0
conv2d_6 (Conv2D)	(None, 14, 14, 96)	153696
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 96)	0
dropout_5 (Dropout)	(None, 7, 7, 96)	0
conv2d_7 (Conv2D)	(None, 7, 7, 128)	307328
max_pooling2d_4 (MaxPooling2D)	(None, 3, 3, 128)	0
dropout_6 (Dropout)	(None, 3, 3, 128)	0
conv2d_8 (Conv2D)	(None, 3, 3, 164)	524964
max_pooling2d_5 (MaxPooling2D)	(None, 1, 1, 164)	0
dropout_7 (Dropout)	(None, 1, 1, 164)	0
flatten_2 (Flatten)	(None, 164)	0
dense_3 (Dense)	(None, 256)	42240
batch_normalization_1 (Batch Normalization)	(None, 256)	1024
dropout_8 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2570
=====		
Total params: 1,083,918		
Trainable params: 1,083,406		
Non-trainable params: 512		
None		

In [12]:

```

convnet2.compile(optimizer=keras.optimizers.Adam(), loss=keras.losses.categorical_crossentropy, metrics=['accuracy'])
convnet2_history=convnet2.fit(x_train,y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test, y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/12
60000/60000 [=====] - 294s 5ms/step - loss: 0.2618 - accuracy: 0.9147 - val_loss: 0.0633 - val_accuracy: 0.9804
Epoch 2/12
60000/60000 [=====] - 290s 5ms/step - loss: 0.0630 - accuracy: 0.9814 - val_loss: 0.0236 - val_accuracy: 0.9920

```

```

al_loss: 0.0186 - val_accuracy: 0.9933
Epoch 4/12
60000/60000 [=====] - 290s 5ms/step - loss: 0.0381 - accuracy: 0.9888 - v
al_loss: 0.0280 - val_accuracy: 0.9924
Epoch 5/12
60000/60000 [=====] - 289s 5ms/step - loss: 0.0342 - accuracy: 0.9896 - v
al_loss: 0.0222 - val_accuracy: 0.9934
Epoch 6/12
60000/60000 [=====] - 290s 5ms/step - loss: 0.0303 - accuracy: 0.9914 - v
al_loss: 0.0280 - val_accuracy: 0.9925
Epoch 7/12
60000/60000 [=====] - 289s 5ms/step - loss: 0.0273 - accuracy: 0.9920 - v
al_loss: 0.0204 - val_accuracy: 0.9945
Epoch 8/12
60000/60000 [=====] - 290s 5ms/step - loss: 0.0244 - accuracy: 0.9930 - v
al_loss: 0.0236 - val_accuracy: 0.9930
Epoch 9/12
60000/60000 [=====] - 291s 5ms/step - loss: 0.0215 - accuracy: 0.9937 - v
al_loss: 0.0196 - val_accuracy: 0.9943
Epoch 10/12
60000/60000 [=====] - 291s 5ms/step - loss: 0.0203 - accuracy: 0.9939 - v
al_loss: 0.0262 - val_accuracy: 0.9932
Epoch 11/12
60000/60000 [=====] - 289s 5ms/step - loss: 0.0204 - accuracy: 0.9941 - v
al_loss: 0.0194 - val_accuracy: 0.9943
Epoch 12/12
60000/60000 [=====] - 291s 5ms/step - loss: 0.0189 - accuracy: 0.9943 - v
al_loss: 0.0209 - val_accuracy: 0.9945

```

In [13]:

```

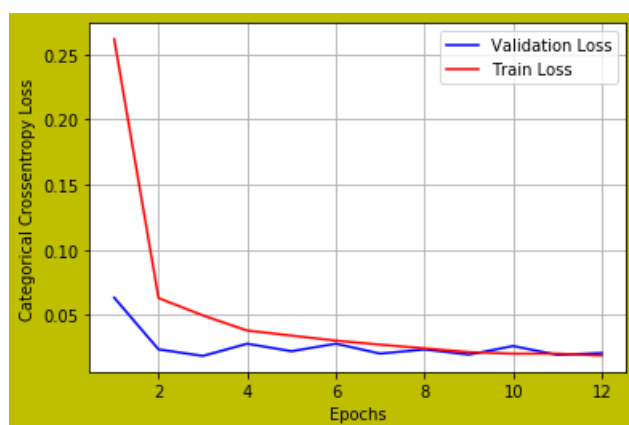
#evaluating model
score=convnet2.evaluate(x_test,y_test,verbose=0)
test_score2=score[0]
test_accuracy2=score[1]
train_accuracy2=max(convnet2_history.history['accuracy'])
print('test score :',test_score2)
print('test Accuracy : ',test_accuracy2)
# error plot
x=list(range(1,epochs+1))
vy=convnet2_history.history['val_loss'] #validation loss
ty=convnet2_history.history['loss'] # train loss
plt_dynamic(x, vy, ty)

```

```

test score : 0.020858116581862943
test Accuracy : 0.9944999814033508

```



3rd Model CNN with 7 ConvNet & 2*2 kernel size :

In [15]:

```

convnet3=Sequential() # Initializing the model
# First ConvNet
convnet3.add(Conv2D(16, kernel_size=(2, 2), activation='relu', padding='same', strides=(1, 1), input_shape=

```

```

convnet3.add(Conv2D(16, kernel_size=(2,2), activation='relu', padding='same', strides=(1,1), input_shape
=input_shape))
convnet3.add(Conv2D(32, kernel_size=(2,2), padding='same', strides=(2,2), activation='relu')) #Second Co
nvnet
convnet3.add(Conv2D(64, kernel_size=(2,2), padding='same', activation='relu')) # 3rd ConvNet
convnet3.add(Dropout(0.15))
convnet3.add(Conv2D(96, kernel_size=(2,2), padding='same', activation='relu')) #fourth Convnet
convnet3.add(MaxPooling2D(pool_size=(2,2)))
convnet3.add(Dropout(0.39))
convnet3.add(Conv2D(128, kernel_size=(2,2), padding='same', activation='relu')) #fifth Convnet
convnet3.add(MaxPooling2D(pool_size=(2,2)))
convnet3.add(Dropout(0.3))
convnet3.add(Conv2D(164, kernel_size=(2,2), padding='same', activation='relu')) #sixth Convnet
convnet3.add(Conv2D(164, kernel_size=(2,2), padding='same', strides=(1,1), activation='relu')) #seventh
Convnet
convnet3.add(MaxPooling2D(pool_size=(2,2)))
convnet3.add(Dropout(0.4))
convnet3.add(Flatten())
#hidden layer
convnet3.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None))) #1 hidden layer
convnet3.add(BatchNormalization())
convnet3.add(Dropout(0.4))
convnet3.add(Dense(148, activation='relu', kernel_initializer=he_normal(seed=None))) #2 hidden layer
convnet3.add(BatchNormalization())
convnet3.add(Dropout(0.3))
convnet3.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None))) #3 hidden layer
convnet3.add(BatchNormalization())
convnet3.add(Dropout(0.4))
convnet3.add(Dense(num_classes, activation='softmax'))
print(convnet3.summary())

```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_16 (Conv2D)	(None, 28, 28, 16)	80
conv2d_17 (Conv2D)	(None, 14, 14, 32)	2080
conv2d_18 (Conv2D)	(None, 14, 14, 64)	8256
dropout_12 (Dropout)	(None, 14, 14, 64)	0
conv2d_19 (Conv2D)	(None, 14, 14, 96)	24672
max_pooling2d_10 (MaxPooling)	(None, 7, 7, 96)	0
dropout_13 (Dropout)	(None, 7, 7, 96)	0
conv2d_20 (Conv2D)	(None, 7, 7, 128)	49280
max_pooling2d_11 (MaxPooling)	(None, 3, 3, 128)	0
dropout_14 (Dropout)	(None, 3, 3, 128)	0
conv2d_21 (Conv2D)	(None, 3, 3, 164)	84132
conv2d_22 (Conv2D)	(None, 3, 3, 164)	107748
max_pooling2d_12 (MaxPooling)	(None, 1, 1, 164)	0
dropout_15 (Dropout)	(None, 1, 1, 164)	0
flatten_3 (Flatten)	(None, 164)	0
dense_5 (Dense)	(None, 256)	42240
batch_normalization_2 (Batch)	(None, 256)	1024
dropout_16 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 148)	38036
batch_normalization_3 (Batch)	(None, 148)	592
dropout_17 (Dropout)	(None, 148)	0

dense_7 (Dense)	(None, 128)	19072
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_18 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290
=====		
Total params: 379,014		
Trainable params: 377,950		
Non-trainable params: 1,064		
None		

In [16]:

```
convnet3.compile(optimizer=keras.optimizers.Adam(),loss=keras.losses.categorical_crossentropy,metrics=['accuracy'])
convnet3_history=convnet3.fit(x_train,y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test, y_test))
```

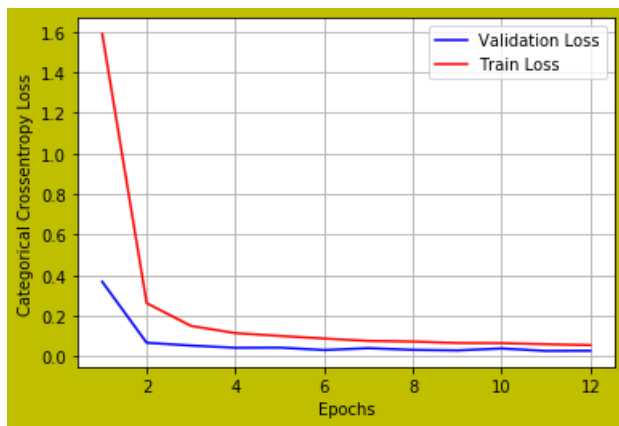
```
Train on 60000 samples, validate on 10000 samples
Epoch 1/12
60000/60000 [=====] - 121s 2ms/step - loss: 1.5887 - accuracy: 0.4850 - val_loss: 0.3675 - val_accuracy: 0.8998
Epoch 2/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.2623 - accuracy: 0.9277 - val_loss: 0.0670 - val_accuracy: 0.9809
Epoch 3/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.1499 - accuracy: 0.9601 - val_loss: 0.0532 - val_accuracy: 0.9863
Epoch 4/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.1145 - accuracy: 0.9696 - val_loss: 0.0424 - val_accuracy: 0.9885
Epoch 5/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.1007 - accuracy: 0.9742 - val_loss: 0.0431 - val_accuracy: 0.9892
Epoch 6/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.0878 - accuracy: 0.9768 - val_loss: 0.0310 - val_accuracy: 0.9908
Epoch 7/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.0761 - accuracy: 0.9794 - val_loss: 0.0404 - val_accuracy: 0.9898
Epoch 8/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.0729 - accuracy: 0.9812 - val_loss: 0.0320 - val_accuracy: 0.9913
Epoch 9/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.0658 - accuracy: 0.9834 - val_loss: 0.0293 - val_accuracy: 0.9922
Epoch 10/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.0651 - accuracy: 0.9829 - val_loss: 0.0388 - val_accuracy: 0.9895
Epoch 11/12
60000/60000 [=====] - 119s 2ms/step - loss: 0.0597 - accuracy: 0.9844 - val_loss: 0.0268 - val_accuracy: 0.9927
Epoch 12/12
60000/60000 [=====] - 118s 2ms/step - loss: 0.0555 - accuracy: 0.9857 - val_loss: 0.0276 - val_accuracy: 0.9928
```

In [17]:

```
score=convnet3.evaluate(x_test,y_test,verbose=0)
test_score3=score[0]
test_accuracy3=score[1]
train_accuracy3=max(convnet3_history.history['accuracy'])
print('test score :',test_score3)
print('test Accuracy :',test_accuracy3)
# error plot
x=list(range(1,epochs+1))
vy=convnet3_history.history['val_loss'] #validation loss
ty=convnet3_history.history['loss'] # train loss
plt_dynamic(x, vy, ty)
```

test_score3 = 0.027557410120102220

test score : 0.027557419139193312
test Accuracy : 0.9927999973297119



Conclusion :

In [18]:

```
from prettytable import PrettyTable
models=['3ConvNet with kernel 3x3',
        '5ConvNet with kernel 5x5',
        '7ConvNet with kernel 2x2']
training_accuracy=[train_accuracy,train_accuracy2,train_accuracy3]
test_accuracy=[test_accuracy,test_accuracy2,test_accuracy3]
INDEX = [1,2,3]
# Initializing prettytable
Model_Performance = PrettyTable()
# Adding columns
Model_Performance.add_column("INDEX.", INDEX)
Model_Performance.add_column("MODEL_NAME",models)
Model_Performance.add_column("TRAINING ACCURACY",training_accuracy)
Model_Performance.add_column("TESTING ACCURACY",test_accuracy)
#Model_Performance.add_column("TEST SCORE",test_score)

# Printing the Model_Performance
print(Model_Performance)
```

INDEX.	MODEL_NAME	TRAINING ACCURACY	TESTING ACCURACY
1	3ConvNet with kernel 3x3	0.9952833	0.9926000237464905
2	5ConvNet with kernel 5x5	0.99431664	0.9944999814033508
3	7ConvNet with kernel 2x2	0.9856667	0.9927999973297119