

Keras -- MLPs on MNIST

In [1]:

```
import matplotlib.pyplot as plt
from keras.utils import np_utils
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
from keras.optimizers import Adam
from keras.layers import Dense, Activation
import seaborn as sns
import numpy as np
import keras
```

Using TensorFlow backend.

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

In [7]:

```
numofclasses = 10
y_train = keras.utils.to_categorical(y_train, numofclasses)
y_test = keras.utils.to_categorical(y_test, numofclasses)
y_train[:5]
```

Out[7]:

```
array([[0., 0., 0., 0., 0., 1., 0., 0., 0., 0.],
       [1., 0., 0., 0., 0., 0., 0., 0., 0., 0.],
       [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],
       [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],
       [0., 0., 0., 0., 0., 0., 0., 0., 0., 1.]], dtype=float32)
```

In [8]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253  70  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

In [9]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
```

```
X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normalizing
print(X_train[0])
```

[illegible]

[illegible]

MODEL 1 WITHOUT BN AND DROPOUT 2 LAYERS :

In [11]:

```
# some model parameters
```

```
output dim = 10
```

```
input_dim = X_train.shape[1]
```

```
batch_size = 128
```

```
nb_epoch = 20
```

In [12]:

```
model1 = Sequential()
model1.add(Dense(364, activation='relu', input_shape=(input_dim, ), kernel_initializer='random_uniform'))
model1.add(Dense(120, activation='relu', kernel_initializer='random_uniform'))
model1.add(Dense(output_dim, activation='softmax'))
print(model1.summary())
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 364)	285740
dense_2 (Dense)	(None, 120)	43800
dense_3 (Dense)	(None, 10)	1210

=====
Total params: 330,750
Trainable params: 330,750
Non-trainable params: 0

None

In [13]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [14]:

```
# Training the model
results = model1.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 57us/step - loss: 0.2798 - accuracy: 0.9214 - val_loss: 0.1341 - val_accuracy: 0.9596

Epoch 2/20

60000/60000 [=====] - 3s 48us/step - loss: 0.1042 - accuracy: 0.9686 - val_loss: 0.0871 - val_accuracy: 0.9723

Epoch 3/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0670 - accuracy: 0.9790 - val_loss: 0.0798 - val_accuracy: 0.9752

Epoch 4/20

60000/60000 [=====] - 3s 51us/step - loss: 0.0471 - accuracy: 0.9855 - val_loss: 0.0749 - val_accuracy: 0.9760

Epoch 5/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0337 - accuracy: 0.9891 - val_loss: 0.0765 - val_accuracy: 0.9769

Epoch 6/20

60000/60000 [=====] - 3s 46us/step - loss: 0.0270 - accuracy: 0.9914 - val_loss: 0.0701 - val_accuracy: 0.9789

Epoch 7/20

60000/60000 [=====] - 3s 46us/step - loss: 0.0202 - accuracy: 0.9939 - val_loss: 0.0808 - val_accuracy: 0.9762

Epoch 8/20

60000/60000 [=====] - 3s 46us/step - loss: 0.0163 - accuracy: 0.9946 - val_loss: 0.0774 - val_accuracy: 0.9784

Epoch 9/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0170 - accuracy: 0.9944 - val_loss: 0.0937 - val_accuracy: 0.9762

Epoch 10/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0137 - accuracy: 0.9953 - val_loss: 0.0723 - val_accuracy: 0.9812

Epoch 11/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0093 - accuracy: 0.9970 - va

```

l_loss: 0.0795 - val_accuracy: 0.9809
Epoch 12/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0104 - accuracy: 0.9967 - va
l_loss: 0.0960 - val_accuracy: 0.9773
Epoch 13/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0119 - accuracy: 0.9961 - va
l_loss: 0.1018 - val_accuracy: 0.9759
Epoch 14/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0097 - accuracy: 0.9969 - va
l_loss: 0.0787 - val_accuracy: 0.9809
Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0077 - accuracy: 0.9977 - va
l_loss: 0.0833 - val_accuracy: 0.9812
Epoch 16/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0087 - accuracy: 0.9967 - va
l_loss: 0.0940 - val_accuracy: 0.9796
Epoch 17/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0064 - accuracy: 0.9977 - va
l_loss: 0.1040 - val_accuracy: 0.9794
Epoch 18/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0073 - accuracy: 0.9975 - va
l_loss: 0.1002 - val_accuracy: 0.9800
Epoch 19/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0095 - accuracy: 0.9970 - va
l_loss: 0.0984 - val_accuracy: 0.9804
Epoch 20/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0053 - accuracy: 0.9984 - va
l_loss: 0.0927 - val_accuracy: 0.9830

```

In [15]:

```
score = model1.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 1s 72us/step
```

In [16]:

```


fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')


# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();

```





In [17]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9829999804496765

2 LAYERS WITH BN AND DROPOUT :

In [18]:

```
model2 = Sequential()

model2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model2.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))
model2.add(Dense(output_dim, activation='softmax'))
print(model2.summary())
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 364)	285740
dense_5 (Dense)	(None, 128)	46720
batch_normalization_1 (Batch Normalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 10)	1290
Total params: 334,262		
Trainable params: 334,006		
Non-trainable params: 256		
None		

In [19]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [20]:

```
results = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

In [21]:

```
score = model2.evaluate(X_test, y_test)
```

10000/10000 [=====] - 0s 44us/step


In [22]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



In [23]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9825000166893005

Changing Dropout to 0.8

In [24]:

```
model2 = Sequential()
model2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model2.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model2.add(BatchNormalization())
model2.add(Dropout(0.8))
model2.add(Dense(output_dim, activation='softmax'))
print(model2.summary())
```

WARNING:tensorflow:Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 364)	285740
dense_8 (Dense)	(None, 128)	46720
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_9 (Dense)	(None, 10)	1290
Total params: 334,262		
Trainable params: 334,006		
Non-trainable params: 256		
None		

In [25]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [26]:

```
results = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 4s 72us/step - loss: 0.4992 - accuracy: 0.8541 - val_loss: 0.1452 - val_accuracy: 0.9573
Epoch 2/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1964 - accuracy: 0.9478 - val_loss: 0.1035 - val_accuracy: 0.9687
Epoch 3/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1378 - accuracy: 0.9629 - val_loss: 0.0976 - val_accuracy: 0.9713
Epoch 4/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1093 - accuracy: 0.9704 - val_loss: 0.0900 - val_accuracy: 0.9725
Epoch 5/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0933 - accuracy: 0.9739 - val_loss: 0.0865 - val_accuracy: 0.9770
Epoch 6/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0737 - accuracy: 0.9795 - val_loss: 0.0785 - val_accuracy: 0.9791
Epoch 7/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0670 - accuracy: 0.9811 - val_loss: 0.0735 - val_accuracy: 0.9802
Epoch 8/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0579 - accuracy: 0.9838 - val_loss: 0.0868 - val_accuracy: 0.9752
Epoch 9/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0513 - accuracy: 0.9852 - val_loss: 0.0798 - val_accuracy: 0.9786
Epoch 10/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0457 - accuracy: 0.9865 - val_loss: 0.0820 - val_accuracy: 0.9792
```

```

1_loss: 0.0929 - val_accuracy: 0.9792
Epoch 11/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0431 - accuracy: 0.9877 - va
l_loss: 0.0830 - val_accuracy: 0.9769
Epoch 12/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0395 - accuracy: 0.9879 - va
l_loss: 0.0916 - val_accuracy: 0.9783
Epoch 13/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0376 - accuracy: 0.9887 - va
l_loss: 0.0997 - val_accuracy: 0.9765
Epoch 14/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0311 - accuracy: 0.9911 - va
l_loss: 0.1028 - val_accuracy: 0.9774
Epoch 15/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0301 - accuracy: 0.9911 - va
l_loss: 0.0919 - val_accuracy: 0.9795
Epoch 16/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0261 - accuracy: 0.9922 - va
l_loss: 0.0871 - val_accuracy: 0.9811
Epoch 17/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0280 - accuracy: 0.9916 - va
l_loss: 0.1085 - val_accuracy: 0.9779
Epoch 18/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0239 - accuracy: 0.9927 - va
l_loss: 0.0917 - val_accuracy: 0.9800
Epoch 19/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0249 - accuracy: 0.9928 - va
l_loss: 0.0943 - val_accuracy: 0.9820
Epoch 20/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0210 - accuracy: 0.9937 - va
l_loss: 0.0877 - val_accuracy: 0.9826

```

In [27]:

```
score = model2.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 0s 44us/step
```

In [28]:

```


fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')


# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();

```





In [29]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9825999736785889

Changing DropOut Rate To 0.2

In [30]:

```
model2 = Sequential()
model2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model2.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model2.add(BatchNormalization())
model2.add(Dropout(0.2))
model2.add(Dense(output_dim, activation='softmax'))
print(model2.summary())
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 364)	285740
dense_11 (Dense)	(None, 128)	46720
batch_normalization_3 (Batch Normalization)	(None, 128)	512
dropout_3 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 10)	1290

Total params: 334,262
Trainable params: 334,006
Non-trainable params: 256

None

In [31]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [32]:

```
results = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 67us/step - loss: 0.2127 - accuracy: 0.9363 - val_loss: 0.1067 - val_accuracy: 0.9690

Epoch 2/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0897 - accuracy: 0.9721 - val_loss: 0.0833 - val_accuracy: 0.9741

Epoch 3/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0615 - accuracy: 0.9812 - val_loss: 0.0780 - val_accuracy: 0.9762

Epoch 4/20

```

60000/60000 [=====] - 3s 55us/step - loss: 0.0446 - accuracy: 0.9855 - va
l_loss: 0.0775 - val_accuracy: 0.9768
Epoch 5/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0361 - accuracy: 0.9886 - va
l_loss: 0.0817 - val_accuracy: 0.9771
Epoch 6/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0288 - accuracy: 0.9906 - va
l_loss: 0.0872 - val_accuracy: 0.9760
Epoch 7/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0243 - accuracy: 0.9920 - va
l_loss: 0.0730 - val_accuracy: 0.9805
Epoch 8/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0232 - accuracy: 0.9927 - va
l_loss: 0.0688 - val_accuracy: 0.9805
Epoch 9/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0205 - accuracy: 0.9932 - va
l_loss: 0.0810 - val_accuracy: 0.9773
Epoch 10/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0158 - accuracy: 0.9949 - va
l_loss: 0.0778 - val_accuracy: 0.9802
Epoch 11/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0143 - accuracy: 0.9952 - va
l_loss: 0.0795 - val_accuracy: 0.9788
Epoch 12/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0151 - accuracy: 0.9949 - va
l_loss: 0.0922 - val_accuracy: 0.9781
Epoch 13/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0132 - accuracy: 0.9956 - va
l_loss: 0.0788 - val_accuracy: 0.9798
Epoch 14/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0107 - accuracy: 0.9962 - va
l_loss: 0.0929 - val_accuracy: 0.9759
Epoch 15/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0148 - accuracy: 0.9949 - va
l_loss: 0.0911 - val_accuracy: 0.9792
Epoch 16/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0107 - accuracy: 0.9964 - va
l_loss: 0.0846 - val_accuracy: 0.9805
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0068 - accuracy: 0.9979 - va
l_loss: 0.0890 - val_accuracy: 0.9790
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0087 - accuracy: 0.9970 - va
l_loss: 0.0911 - val_accuracy: 0.9795
Epoch 19/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0111 - accuracy: 0.9963 - va
l_loss: 0.0888 - val_accuracy: 0.9774
Epoch 20/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0099 - accuracy: 0.9966 - va
l_loss: 0.0831 - val_accuracy: 0.9816

```

In [33]:

```
score = model2.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 0s 46us/step
```

In [34]:

```


fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();

```



In [35]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9815999865531921

Model With 3 Hidden Layers :

In [36]:

```
model3 = Sequential()
model3.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model3.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model3.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model3.add(Dense(output_dim, activation='softmax'))
print(model3.summary())
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 364)	285740
dense_14 (Dense)	(None, 128)	46720
dense_15 (Dense)	(None, 64)	8256
dense_16 (Dense)	(None, 10)	650

Total params: 341,366
Trainable params: 341,366
Non-trainable params: 0

None

In [37]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [38]:

```
results = model3.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20


In [39]:

```
score = model3.evaluate(X_test, y_test)
```

10000/10000 [=====] - 0s 50us/step

In [40]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')
# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))
train_loss = results.history['loss']
val_loss = results.history['val_loss']
ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



2.5 5.0 7.5 10.0 12.5 15.0 17.5 20.0
Epochs

In [41]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9761999845504761

Model with 3 hidden layers with BN & DropOut :

In [86]:

```
model3 = Sequential()  
model3.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))  
model3.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))  
model3.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))  
model3.add(BatchNormalization())  
model3.add(Dropout(0.5))  
model3.add(Dense(output_dim, activation='softmax'))  
print(model3.summary())
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
=====		
dense_53 (Dense)	(None, 364)	285740

dense_54 (Dense)	(None, 128)	46720

dense_55 (Dense)	(None, 64)	8256

batch_normalization_10 (Batch Normalization)	(None, 64)	256

dropout_10 (Dropout)	(None, 64)	0

dense_56 (Dense)	(None, 10)	650
=====		
Total params: 341,622		
Trainable params: 341,494		
Non-trainable params: 128		

None		

In [87]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [88]:

```
results = model3.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20

```
Epoch 14/20  
Epoch 15/20  
Epoch 16/20  
Epoch 17/20  
Epoch 18/20  
Epoch 19/20  
Epoch 20/20
```


In [89]:

```
score = model3.evaluate(X_test, y_test)
```

10000/10000 [=====] - 1s 54us/step

In [90]:

```
fig,ax = plt.subplots(1,1)  
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')  
# list of epoch numbers  
list_of_epoch = list(range(1,nb_epoch+1))  
train_loss = results.history['loss']  
val_loss = results.history['val_loss']  
ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")  
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")  
plt.legend()  
plt.grid()  
plt.show();
```



In [91]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9818000197410583

DropOut with 0.8

In [92]:

```
model4 = Sequential()  
model4.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform',  
...))
```



```
m'))
model4.add(Dense(128, activation='relu',kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model4.add(BatchNormalization())
model4.add(Dropout(0.8))
model4.add(Dense(output_dim, activation='softmax'))
print(model4.summary())
```

WARNING:tensorflow:Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_57 (Dense)	(None, 364)	285740
dense_58 (Dense)	(None, 128)	46720
dense_59 (Dense)	(None, 64)	8256
batch_normalization_11 (Batch Normalization)	(None, 64)	256
dropout_11 (Dropout)	(None, 64)	0
dense_60 (Dense)	(None, 10)	650

Total params: 341,622
 Trainable params: 341,494
 Non-trainable params: 128

None

In [93]:

```
model4.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accuracy'])
```

In [94]:

```
results = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 91us/step - loss: 0.5065 - accuracy: 0.8498 - val_loss: 0.1633 - val_accuracy: 0.9533

Epoch 2/20

60000/60000 [=====] - 4s 69us/step - loss: 0.2083 - accuracy: 0.9464 - val_loss: 0.1047 - val_accuracy: 0.9675

Epoch 3/20

60000/60000 [=====] - 4s 72us/step - loss: 0.1507 - accuracy: 0.9619 - val_loss: 0.1190 - val_accuracy: 0.9679

Epoch 4/20

60000/60000 [=====] - 4s 65us/step - loss: 0.1111 - accuracy: 0.9704 - val_loss: 0.0968 - val_accuracy: 0.9742

Epoch 5/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0969 - accuracy: 0.9743 - val_loss: 0.1000 - val_accuracy: 0.9737

Epoch 6/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0843 - accuracy: 0.9776 - val_loss: 0.1257 - val_accuracy: 0.9717

Epoch 7/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0712 - accuracy: 0.9806 - val_loss: 0.0896 - val_accuracy: 0.9783

Epoch 8/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0629 - accuracy: 0.9821 - val_loss: 0.1210 - val_accuracy: 0.9756

Epoch 9/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0592 - accuracy: 0.9829 - val_loss: 0.0984 - val_accuracy: 0.9795

Epoch 10/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0490 - accuracy: 0.9848 - val_loss: 0.1216 - val_accuracy: 0.9752

Epoch 11/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0461 - accuracy: 0.9860 - va

```

l_loss: 0.1012 - val_accuracy: 0.9801
Epoch 12/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0478 - accuracy: 0.9847 - va
l_loss: 0.1217 - val_accuracy: 0.9784
Epoch 13/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0389 - accuracy: 0.9877 - va
l_loss: 0.1106 - val_accuracy: 0.9817
Epoch 14/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0416 - accuracy: 0.9879 - va
l_loss: 0.1341 - val_accuracy: 0.9790
Epoch 15/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0392 - accuracy: 0.9873 - va
l_loss: 0.1147 - val_accuracy: 0.9820
Epoch 16/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0382 - accuracy: 0.9869 - va
l_loss: 0.1445 - val_accuracy: 0.9782
Epoch 17/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0303 - accuracy: 0.9893 - va
l_loss: 0.1444 - val_accuracy: 0.9773
Epoch 18/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0385 - accuracy: 0.9869 - va
l_loss: 0.1094 - val_accuracy: 0.9820
Epoch 19/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0268 - accuracy: 0.9903 - va
l_loss: 0.1085 - val_accuracy: 0.9825
Epoch 20/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0267 - accuracy: 0.9906 - va
l_loss: 0.1142 - val_accuracy: 0.9813

```

In [95]:

```
score = model4.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 1s 56us/step
```

In [96]:

```


fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')


# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();

```





In [97]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9812999963760376

changing dropout rate to 0.2

In [55]:

```
model4 = Sequential()
model4.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model4.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model4.add(BatchNormalization())
model4.add(Dropout(0.2))
model4.add(Dense(output_dim, activation='softmax'))
print(model4.summary())
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_25 (Dense)	(None, 364)	285740
dense_26 (Dense)	(None, 128)	46720
dense_27 (Dense)	(None, 64)	8256
batch_normalization_6 (Batch Normalization)	(None, 64)	256
dropout_6 (Dropout)	(None, 64)	0
dense_28 (Dense)	(None, 10)	650
Total params: 341,622		
Trainable params: 341,494		
Non-trainable params: 128		

None

In [56]:

```
model4.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [57]:

```
results = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 77us/step - loss: 0.2278 - accuracy: 0.9344 - val_loss: 0.1169 - val_accuracy: 0.9688

Epoch 2/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0891 - accuracy: 0.9729 - val_loss: 0.1046 - val_accuracy: 0.9671

Epoch 3/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0589 - accuracy: 0.9817 - val_loss: 0.1284 - val_accuracy: 0.9610

Epoch 4/20

```

Epoch 4/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0467 - accuracy: 0.9857 - va
l_loss: 0.0777 - val_accuracy: 0.9768
Epoch 5/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0368 - accuracy: 0.9892 - va
l_loss: 0.0656 - val_accuracy: 0.9808
Epoch 6/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0267 - accuracy: 0.9916 - va
l_loss: 0.0695 - val_accuracy: 0.9808
Epoch 7/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0252 - accuracy: 0.9920 - va
l_loss: 0.0647 - val_accuracy: 0.9817
Epoch 8/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0198 - accuracy: 0.9937 - va
l_loss: 0.0855 - val_accuracy: 0.9775
Epoch 9/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0182 - accuracy: 0.9942 - va
l_loss: 0.0880 - val_accuracy: 0.9767
Epoch 10/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0188 - accuracy: 0.9941 - va
l_loss: 0.0762 - val_accuracy: 0.9806
Epoch 11/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0134 - accuracy: 0.9957 - va
l_loss: 0.0832 - val_accuracy: 0.9790
Epoch 12/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0137 - accuracy: 0.9958 - va
l_loss: 0.0663 - val_accuracy: 0.9829
Epoch 13/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0092 - accuracy: 0.9972 - va
l_loss: 0.0887 - val_accuracy: 0.9787
Epoch 14/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0156 - accuracy: 0.9951 - va
l_loss: 0.0791 - val_accuracy: 0.9816
Epoch 15/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0101 - accuracy: 0.9969 - va
l_loss: 0.0879 - val_accuracy: 0.9806
Epoch 16/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0112 - accuracy: 0.9965 - va
l_loss: 0.0699 - val_accuracy: 0.9821
Epoch 17/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0095 - accuracy: 0.9970 - va
l_loss: 0.0826 - val_accuracy: 0.9814
Epoch 18/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0063 - accuracy: 0.9979 - va
l_loss: 0.0756 - val_accuracy: 0.9843
Epoch 19/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0093 - accuracy: 0.9973 - va
l_loss: 0.0848 - val_accuracy: 0.9791
Epoch 20/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0067 - accuracy: 0.9979 - va
l_loss: 0.0775 - val_accuracy: 0.9838

```

In [58]:

```
score = model4.evaluate(X_test, y_test)
```


```
10000/10000 [=====] - 1s 55us/step
```

In [59]:

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')
list_of_epoch = list(range(1,nb_epoch+1))
train_loss = results.history['loss']
val_loss = results.history['val_loss']
ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();

```



In [60]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9837999939918518

Model with 5 hidden layers :

In [61]:

```
model5 = Sequential()
model5.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model5.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(output_dim, activation='softmax'))
print(model5.summary())
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 364)	285740
dense_30 (Dense)	(None, 128)	46720
dense_31 (Dense)	(None, 96)	12384
dense_32 (Dense)	(None, 64)	6208
dense_33 (Dense)	(None, 32)	2080
dense_34 (Dense)	(None, 10)	330
Total params: 353,462		
Trainable params: 353,462		
Non-trainable params: 0		

None

In [62]:

```
model5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model5.compile(optimizer='adam',loss='categorical_crossentropy',metrics=['accuracy'])
```

In [63]:

```
results = model5.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
Epoch 2/20
Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

In [64]:

```
score = model5.evaluate(X_test, y_test)
```

10000/10000 [=====] - 1s 63us/step


In [65]:


```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```





In [66]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9768999814987183

Model With 5 hidden layers along with BN and Dropout

In [67]:

```
model6 = Sequential()
model6.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model6.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))
model6.add(BatchNormalization())
model6.add(Dropout(0.5))
model6.add(Dense(output_dim, activation='softmax'))
print(model6.summary())
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 364)	285740
dense_36 (Dense)	(None, 128)	46720
dense_37 (Dense)	(None, 96)	12384
dense_38 (Dense)	(None, 64)	6208
dense_39 (Dense)	(None, 32)	2080
batch_normalization_7 (Batch Normalization)	(None, 32)	128
dropout_7 (Dropout)	(None, 32)	0
dense_40 (Dense)	(None, 10)	330
Total params: 353,590		
Trainable params: 353,526		
Non-trainable params: 64		
None		

In [68]:

```
model6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [69]:

```
results = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
Epoch 2/20

Epoch 3/20
Epoch 4/20
Epoch 5/20
Epoch 6/20
Epoch 7/20
Epoch 8/20
Epoch 9/20
Epoch 10/20
Epoch 11/20
Epoch 12/20
Epoch 13/20
Epoch 14/20
Epoch 15/20
Epoch 16/20
Epoch 17/20
Epoch 18/20
Epoch 19/20
Epoch 20/20

In [70]:

```
score = model6.evaluate(X_test, y_test)
```

10000/10000 [=====] - 1s 56us/step


In [71]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



In [72]:


```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9768000245094299

Dropout to 0.8

In [73]:

```
model6 = Sequential()
model6.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model6.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))
model6.add(BatchNormalization())
model6.add(Dropout(0.8))
model6.add(Dense(output_dim, activation='softmax'))
print(model6.summary())
```

WARNING:tensorflow:Large dropout rate: 0.8 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Model: "sequential_11"

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 364)	285740
dense_42 (Dense)	(None, 128)	46720
dense_43 (Dense)	(None, 96)	12384
dense_44 (Dense)	(None, 64)	6208
dense_45 (Dense)	(None, 32)	2080
batch_normalization_8 (Batch Normalization)	(None, 32)	128
dropout_8 (Dropout)	(None, 32)	0
dense_46 (Dense)	(None, 10)	330
Total params: 353,590		
Trainable params: 353,526		
Non-trainable params: 64		
None		

In [74]:

```
model6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [75]:

```
results = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 86us/step - loss: 0.8172 - accuracy: 0.7139 - val_loss: 0.2154 - val_accuracy: 0.9497

Epoch 2/20

60000/60000 [=====] - 4s 65us/step - loss: 0.4396 - accuracy: 0.8383 - val_loss: 0.1347 - val_accuracy: 0.9660

Epoch 3/20

60000/60000 [=====] - 4s 62us/step - loss: 0.3578 - accuracy: 0.8597 - val_loss: 0.1160 - val_accuracy: 0.9722

Epoch 4/20

60000/60000 [=====] - 4s 62us/step - loss: 0.3197 - accuracy: 0.8700 - val_loss: 0.0984 - val_accuracy: 0.9756

```

1_loss: 0.0904 - val_accuracy: 0.9750
Epoch 5/20
60000/60000 [=====] - 4s 62us/step - loss: 0.3013 - accuracy: 0.8733 - va
1_loss: 0.0975 - val_accuracy: 0.9775
Epoch 6/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2850 - accuracy: 0.8792 - va
1_loss: 0.1093 - val_accuracy: 0.9766
Epoch 7/20
60000/60000 [=====] - 4s 63us/step - loss: 0.2698 - accuracy: 0.8864 - va
1_loss: 0.1413 - val_accuracy: 0.9700
Epoch 8/20
60000/60000 [=====] - 4s 64us/step - loss: 0.2617 - accuracy: 0.8892 - va
1_loss: 0.1271 - val_accuracy: 0.9759
Epoch 9/20
60000/60000 [=====] - 4s 62us/step - loss: 0.2481 - accuracy: 0.8956 - va
1_loss: 0.1172 - val_accuracy: 0.9761
Epoch 10/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2380 - accuracy: 0.8997 - va
1_loss: 0.1037 - val_accuracy: 0.9792
Epoch 11/20
60000/60000 [=====] - 4s 64us/step - loss: 0.2433 - accuracy: 0.8986 - va
1_loss: 0.1231 - val_accuracy: 0.9787
Epoch 12/20
60000/60000 [=====] - 4s 61us/step - loss: 0.2241 - accuracy: 0.9064 - va
1_loss: 0.1119 - val_accuracy: 0.9817
Epoch 13/20
60000/60000 [=====] - 4s 60us/step - loss: 0.2055 - accuracy: 0.9132 - va
1_loss: 0.1059 - val_accuracy: 0.9821
Epoch 14/20
60000/60000 [=====] - 4s 63us/step - loss: 0.2039 - accuracy: 0.9183 - va
1_loss: 0.1331 - val_accuracy: 0.9791
Epoch 15/20
60000/60000 [=====] - 4s 63us/step - loss: 0.1944 - accuracy: 0.9209 - va
1_loss: 0.1355 - val_accuracy: 0.9786
Epoch 16/20
60000/60000 [=====] - 4s 63us/step - loss: 0.1909 - accuracy: 0.9203 - va
1_loss: 0.1303 - val_accuracy: 0.9818
Epoch 17/20
60000/60000 [=====] - 4s 64us/step - loss: 0.1806 - accuracy: 0.9239 - va
1_loss: 0.1263 - val_accuracy: 0.9824
Epoch 18/20
60000/60000 [=====] - 4s 63us/step - loss: 0.1823 - accuracy: 0.9261 - va
1_loss: 0.1418 - val_accuracy: 0.9819
Epoch 19/20
60000/60000 [=====] - 4s 64us/step - loss: 0.1870 - accuracy: 0.9252 - va
1_loss: 0.1468 - val_accuracy: 0.9812
Epoch 20/20
60000/60000 [=====] - 4s 64us/step - loss: 0.1823 - accuracy: 0.9261 - va
1_loss: 0.1471 - val_accuracy: 0.9811

```

In [76]:

```
score = model6.evaluate(X_test, y_test)
```

```
10000/10000 [=====] - 1s 56us/step
```

In [77]:

```


fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();

```



In [78]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9811000227928162

DropOut to 0.2

In [79]:

```
model6 = Sequential()
model6.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model6.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))
model6.add(BatchNormalization())
model6.add(Dropout(0.2))
model6.add(Dense(output_dim, activation='softmax'))
print(model6.summary())
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 364)	285740
dense_48 (Dense)	(None, 128)	46720
dense_49 (Dense)	(None, 96)	12384
dense_50 (Dense)	(None, 64)	6208
dense_51 (Dense)	(None, 32)	2080
batch_normalization_9 (Batch Normalization)	(None, 32)	128
dropout_9 (Dropout)	(None, 32)	0
dense_52 (Dense)	(None, 10)	330

Total params: 353,590

Trainable params: 353,590

trainable params: 555,520
Non-trainable params: 64

None

In [80]:

```
model6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [81]:

```
results = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 88us/step - loss: 0.3328 - accuracy: 0.9092 - val_loss: 0.1275 - val_accuracy: 0.9672
Epoch 2/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1045 - accuracy: 0.9708 - val_loss: 0.0985 - val_accuracy: 0.9708
Epoch 3/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0696 - accuracy: 0.9801 - val_loss: 0.1003 - val_accuracy: 0.9707
Epoch 4/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0529 - accuracy: 0.9850 - val_loss: 0.0821 - val_accuracy: 0.9755
Epoch 5/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0427 - accuracy: 0.9882 - val_loss: 0.0860 - val_accuracy: 0.9768
Epoch 6/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0341 - accuracy: 0.9898 - val_loss: 0.0892 - val_accuracy: 0.9767
Epoch 7/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0269 - accuracy: 0.9923 - val_loss: 0.0939 - val_accuracy: 0.9757
Epoch 8/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0264 - accuracy: 0.9919 - val_loss: 0.0809 - val_accuracy: 0.9806
Epoch 9/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0229 - accuracy: 0.9932 - val_loss: 0.0822 - val_accuracy: 0.9785
Epoch 10/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0192 - accuracy: 0.9944 - val_loss: 0.0980 - val_accuracy: 0.9785
Epoch 11/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0191 - accuracy: 0.9942 - val_loss: 0.0891 - val_accuracy: 0.9797
Epoch 12/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0164 - accuracy: 0.9954 - val_loss: 0.0855 - val_accuracy: 0.9800
Epoch 13/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0143 - accuracy: 0.9957 - val_loss: 0.0841 - val_accuracy: 0.9807
Epoch 14/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0161 - accuracy: 0.9954 - val_loss: 0.0692 - val_accuracy: 0.9836
Epoch 15/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0127 - accuracy: 0.9959 - val_loss: 0.0786 - val_accuracy: 0.9819
Epoch 16/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0136 - accuracy: 0.9958 - val_loss: 0.0804 - val_accuracy: 0.9813
Epoch 17/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0119 - accuracy: 0.9965 - val_loss: 0.0800 - val_accuracy: 0.9823
Epoch 18/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0112 - accuracy: 0.9968 - val_loss: 0.0887 - val_accuracy: 0.9810
Epoch 19/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0074 - accuracy: 0.9978 - val_loss: 0.1007 - val_accuracy: 0.9787
Epoch 20/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0137 - accuracy: 0.9961 - val_loss: 0.0786 - val_accuracy: 0.9826
```

In [82]:

```
score = model6.evaluate(X_test, y_test)
```

10000/10000 [=====] - 1s 55us/step


In [83]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = results.history['loss']
val_loss = results.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



In [84]:

```
print('Test accuracy:', score[1])
```

Test accuracy: 0.9825999736785889

Conclusions

In [85]:

```
# Please compare all your models using Prettytable library
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Numer of Layers", "BN", "Dropout", "Accuracy"]
x.add_row(["2", "NO", "NO", 0.981])
x.add_row(["2", "YES", 0.5, 0.982])
```

```

x.add_row(["2", 'YES', 0.5, 0.982])
x.add_row(["2", 'YES', 0.8, 0.982])
x.add_row(["2", 'YES', 0.2, 0.981])
x.add_row(["3", 'NO', "NO", 0.971])
x.add_row(["3", 'YES', 0.5, 0.976])
x.add_row(["3", 'YES', 0.8, 0.980])
x.add_row(["3", 'YES', 0.2, 0.983])
x.add_row(["5", 'NO', 'NO', 0.976])
x.add_row(["5", 'YES', 0.5, 0.976])
x.add_row(["5", 'YES', 0.8, 0.981])
x.add_row(["5", 'YES', 0.2, 0.982])
print(x)

```

Numer of Layers	BN	Dropout	Accuracy
2	NO	NO	0.981
2	YES	0.5	0.982
2	YES	0.8	0.982
2	YES	0.2	0.981
3	NO	NO	0.971
3	YES	0.5	0.976
3	YES	0.8	0.98
3	YES	0.2	0.983
5	NO	NO	0.976
5	YES	0.5	0.976
5	YES	0.8	0.981
5	YES	0.2	0.982